


Gather



Clean



Model



Visualize



Data Modeling for Experts

It is time to take the next step with your data modeling skills!

kapacity Fellowwind



NEXTAGENDA»
ACHIEVE MORE WITH YOUR DATA
DIGITAL WORKPLACE SOLUTIONS
FROM KONICA MINOLTA



Initiate



RELATE IT

[stoltze][it]unit^{it}



Marc Lelijveld

Data & Analytics consultant
Macaw Netherlands



Marc.Lelijveld@outlook.com



[@MarcLelijveld](https://twitter.com/MarcLelijveld)



linkedin.com/in/MarcLelijveld



Data-Marc.com



After this session you will be able to:



Successfully implement
calculation groups



Build and optimize
composite models




Combine DirectQuery and
import storage modes in a
single
data model



Optimize query
performance in complex
situations

What are we talking about?

Gather



Clean



Model





Visualize



Assumption of your basic understanding

- The importance of a star schema
- Different data model types
- Modeling best practices, including naming conventions for tables and columns
- Auto date/time
- Relationship direction
- Limited and regular relationships
- Ambiguous data models
- Aggregations
- Dealing with multiple fact tables
- Different types of fact tables
- Slowly changing dimensions
- Data groups and hierarchies



Storage modes



Different types of storage modes

Three storage modes

- **Import** – data cached in the model
- **DirectQuery** – queries are submitted to the back-end data source
- **Dual** – can act in both above storage modes, depending on query context

Configuring storage modes

- Storage modes are set on table level
- Setting storage mode to Import is an irreversible operation
- Data in DirectQuery mode cannot be displayed in the data tab

The screenshot displays the Power BI Desktop interface. The ribbon at the top includes tabs for File, Home, Help, and External Tools. The Home tab is active, showing various tool groups like Clipboard, Data, Queries, Relationships, Security, and Q&A. The main workspace shows a data model with four tables: Product Subcategory, Product, Internet Sales, and Internet Sales - Agg. The Internet Sales table is highlighted with a yellow border. The Properties pane on the right shows the 'Advanced' tab with the 'Storage mode' set to 'DirectQuery'. The 'Internet Sales' table is also highlighted with a yellow border in the model view.

Understanding Dual mode

Example

- Internet Sales → DirectQuery
- Internet Sales - Agg → Import
- Product Category → Dual

The screenshot displays the Microsoft Power BI Desktop interface. The top ribbon includes tabs for File, Home, Help, and External Tools. The Home tab is active, showing various toolbars for Clipboard, Data, Queries, Relationships, Security, and Q&A. The main workspace shows a data model with four tables: Product Subcategory, Product, Internet Sales, and Internet Sales - Agg. The tables are connected with relationships. The 'Internet Sales' table is highlighted with a yellow border. The 'Properties' pane on the right shows the 'Advanced' section with 'Storage mode' set to 'DirectQuery'.

Product Subcategory

- English Product Subcategory N...
- French Product Subcategory Na...
- ProductCategoryKey
- ProductSubcategoryAlt...
- ProductSubcategoryKey
- Spanish Product Subcategory N...

Product

- Arabic Description
- Chinese Description
- Class
- Color
- Days To Manufacture
- Dealer Price

Internet Sales

- Carrier Tracking Number
- CurrencyKey
- Customer PO Number
- CustomerKey
- Discount Amount
- DueDateKey
- Extended Amount
- Freight Amount
- Order Calendar Year

Internet Sales - Agg

- Count
- Order Calendar Year
- ProductSubcategoryKey
- Sales Amount

Properties

General

Name: Internet Sales

Description: Enter a description

Synonyms: Enter a comma-separated list of synonyms for Q&A

Row label: Select a row label

Key column: Select a column with unique values

Is hidden: No

Is featured table: No

Advanced

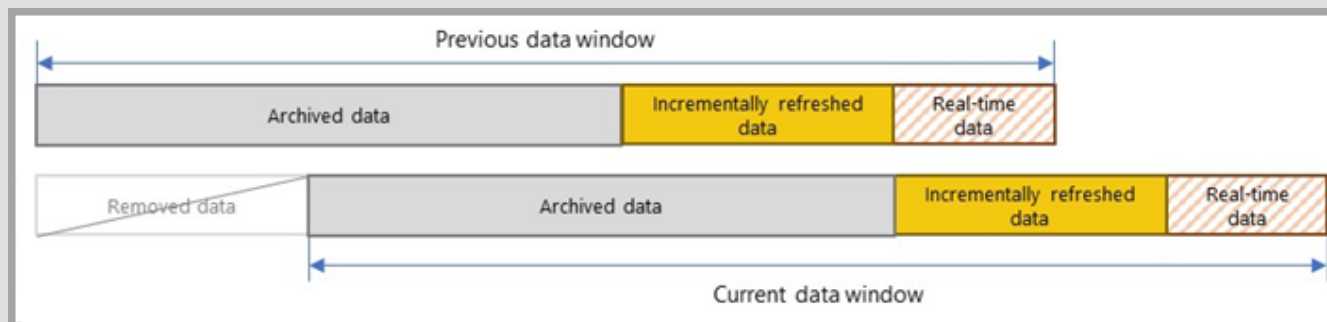
Storage mode: DirectQuery

Hybrid tables

There is one exception...

- Storage mode is set on partition level.
- Only possible using the Power BI interface for Hybrid Tables, not possible to manually configure.
- Hybrid tables goes hand in hand with incremental refresh.

Granularity	Name	Row Count	
Year	2011	295,489,717	Archived: Import
Year	2012	297,678,498	
Year	2013	295,575,442	
Year	2014	292,477,875	
Year	2015	297,780,469	
Year	2016	294,060,081	
Year	2017	300,419,682	
Year	2018	296,541,108	
Year	2019	292,787,420	
Year	2020	299,273,979	
Quarter	2021Q1	74,135,277	Incremental refresh: Import
Month	2021Q104	24,939,498	
Day	2021Q10501	820,805	
Day	2021Q10502	826,885	Real time: DirectQuery
Day	2021Q10503	821,043	
Day-DirectQuery	2021Q10504-DQ	271,110	
Total		3,063,898,887	



Benefits

Benefits to choose one or the other

- Improve query performance – cache data for faster end-user performance
- Data refresh optimization – no need to refresh for non-cached data
- Near-real time requirements – reduce query latency when in DirectQuery mode
- Large datasets – choose to not import certain data



Caches and DirectQuery

Risks of mixing storage modes

- Avoid mismatch in results when query bits DirectQuery compared to Import
- Data cached (import) could be behind compared to DirectQuery data
- Make sure cached data is kept **in sync** – regularly refresh!



Behavior of Dual mode

Query context defines storage mode used

- Sales per product subcategory, both tables will use Import mode
- Sales per product, both tables will hit DirectQuery
- Extended Amount per product subcategory, both tables will hit DirectQuery

The screenshot displays the Microsoft Power BI Desktop interface. The main workspace shows a data model with three tables: 'Product Subcategory', 'Product', and 'Internet Sales'. The 'Product Subcategory' table is connected to the 'Product' table, and the 'Product' table is connected to the 'Internet Sales' table. The 'Internet Sales' table is highlighted with a yellow border. The 'Properties' pane on the right shows the 'Storage mode' set to 'DirectQuery'.

Product Subcategory Table:

- English Product Subcategory N...
- French Product Subcategory Na...
- ProductCategoryKey
- ProductSubcategoryAlt...
- ProductSubcategoryKey
- Spanish Product Subcategory N...

Product Table:

- Arabic Description
- Chinese Description
- Class
- Color
- Days To Manufacture
- Dealer Price

Internet Sales Table:

- Carrier Tracking Number
- CurrencyKey
- Customer PO Number
- CustomerKey
- Discount Amount
- DueDateKey
- Extended Amount
- Freight Amount
- Order Calendar Year

Properties Pane:

- Name: Internet Sales
- Description: Enter a description
- Synonyms: Enter a comma-separated list of synonyms for Q&A
- Row label: Select a row label
- Key column: Select a column with unique values
- Is hidden: No
- Is featured table: No
- Storage mode: DirectQuery

Composite models



Introducing composite models


Composite models do either or both of the following actions

- Combine data from one or more DirectQuery sources
- Combines data from DirectQuery and import data

Three related features in Power BI

- Composite models, to allow different connections and storage models in a single data model
- Many-to-many relationships
- Storage modes

Limited relationships might occur, when relationships cross source group



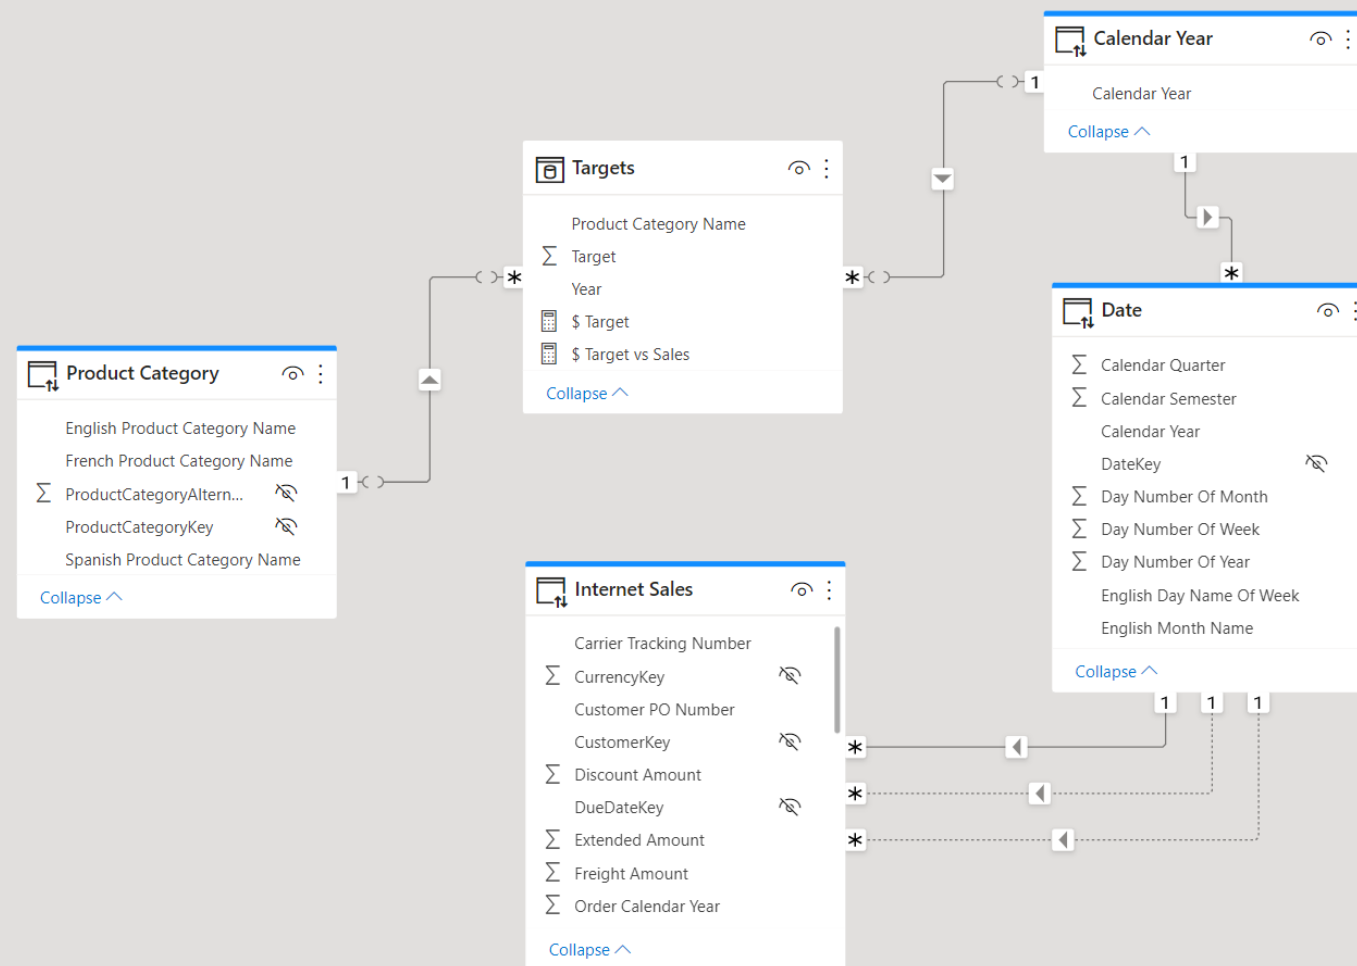
Example, sales model with local target sheet

Example use case for composite models in Power BI

- Centrally maintained data model, containing all sales
- Region has their own sheet with targets in Excel, which they can blend with the central data model

Source group – everything loaded from one DirectQuery source

All import tables and calculated tables are in one source group

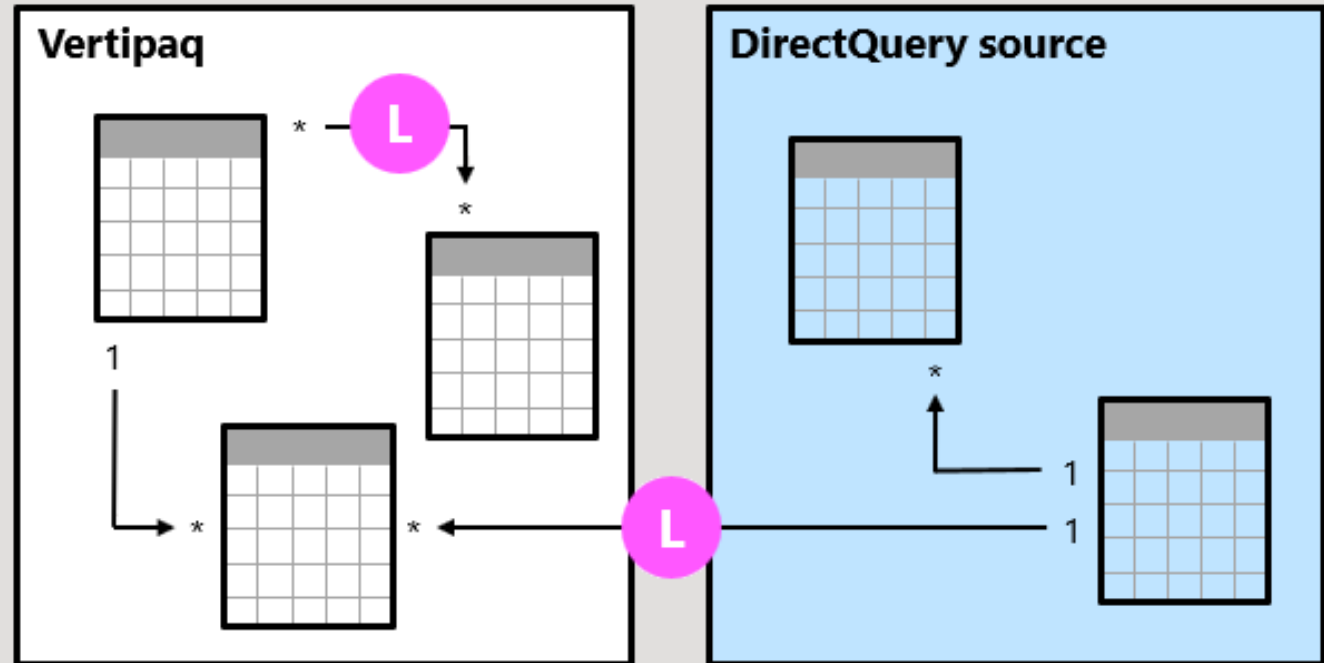


Limited relationships revisited

A model relationship is *limited* when there's no guaranteed "one" side

It can be the case for two reasons:

- The relationship uses a Many-to-many cardinality type (even if one or both columns contain unique values)
- The relationship is cross source group (which can only ever be the case for Composite models)
- It can be a cross-storage mode but does not have to be



Composite models

Demo

Security implications

Composite models do either or both of the following actions

- A query sent to one data source can include data values that have been retrieved from another data source

Potentially

- Any administrator of the data sources, who can view the traces or audit logs could view this information, without permissions to the original data source
- The encryption setting for each source should be considered, to avoid breaching encryption while the data is included in a query sent to another unencrypted source

Potential security risk

When you add a second data source, information from one data source might be included in queries that are sent to other data sources. Only add another source if you trust the owners and admins with the data.

Do you want to build this model?

[Learn more](#)

OK

Cancel

```
SELECT ...  
FROM ...  
inner join (  
  (SELECT N'Andersen' AS [c52],N'Mountain-200 Black, 38' AS [c4] ) UNION ALL  
  (SELECT N'Andersen' AS [c52],N'Mountain-200 Black, 42' AS [c4] ) UNION ALL  
  (SELECT N'Bell' AS [c52],N'Mountain-200 Black, 46' AS [c4] ) UNION ALL  
  (SELECT N'Bell' AS [c52],N'Mountain-500 Black, 52' AS [c4] ) UNION ALL  
  ...
```



Performance implications

Composite models might have slower performance

- **General limitations of DirectQuery** apply, make sure the back-end data source has sufficient resources to provide best performance as possible
- Query execution potentially needs to **cross sources**, where values from one source needs to be evaluated in another source (rich set of values in the WHERE clause)
- **Difference in granularity** of data across sources, where on-demand aggregation is required before displaying data in a visual
- **Multiple queries** must be executed to retrieve results (typically with DistinctCount)



Composite models using Power BI datasets and Azure Analysis Services

With the ability to connect to Power BI datasets and Azure Analysis Services models using DirectQuery (in preview since December 2020) you can now build a composite model using these sources.

Note the following design goals / non-goals for this feature:

Goals:

- **Extend:** Add data to enterprise semantic model
- **Enrich:** Add measure, calculation, formatting to enrich enterprise semantic model without adding any data
- **Overview:** disjointed datasets to give business overview ('dashboard')

Non goals:

- Combine two or more sources by joining large shared dimension tables
- Combine large dimension table and large fact table from different sources

Composite models 2

Demo

Calculation groups



Introducing calculation groups

Benefits

- Reduce the number of redundant measures and grouping common measure expressions as calculation items
- Avoids duplicating logic in different measures
- Typical use cases are
 - Time-intelligence calculations (YTD / QTD / MTD / ...)
 - Format string change, like currency conversions

Limitations

- Can only be created from external tools in Power BI (Any tool using the XMLA endpoint such as Tabular Editor)
- Object level security on Calculation group items is not supported
- Smart narrative visuals in Power BI are not supported with Calculation Groups
- Be aware of potential unexpected behavior in combination with Composite Models



How calculation groups work

Specific DAX expressions for Calculation Groups

- SELECTEDMEASURE()
- SELECTEDMEASURENAME()
- ISSELECTEDMEASURE()
- SELECTEDMEASUREFORMATSTRING()

Classic measure:

```
MTD =  
CALCULATE (  
    SUM ( Sales[SalesAmount] ),  
    DATESMTD ( DimDate[Date] )  
)
```

Measure with MTD context with Calculation Group:

```
MTD =  
CALCULATE (  
    SELECTEDMEASURE (),  
    DATESMTD ( DimDate[Date] )  
)
```

Create and use calculation groups

Demo

Wholesale and non-wholesale queries



Remote vs local model

Remote model – a model that serves as a data source for the current (local) model

- Remote / vs local table – a remote table is a table that is sourced from a remote model. Any table that is not from a remote model is considered local
- Remote measures are treated as a black box to the composite model: a calc groups might be applied, or not, depends on wholesale vs non-wholesale



Wholesale vs non-wholesale

- **Wholesale:** depends on objects in single remote source group
 - Calculated columns defined on a table from a remote source group must be wholesale
- **Non-wholesale:** depends on objects in more than one source group



Calculation groups and measures in a composite model

Depending on the configuration, calculation groups and measures have interesting interactions

Measure location	Is measure wholesaleable	Remote calc group applied
Remote	Yes	Yes
Local	Yes	Yes
Local	No (either because it depends on more than one source group or just the local source group)	No – but can be indirectly

Remember: measures might be impacted by remote calculation groups, when applied

Calculation groups and wholesale / non-wholesale measures

Demo

Key take-aways

- Consider picking the right storage mode for your need
- Be aware of the security and performance implications of composite models
- Reduce the number of redundant measures and avoid duplication calculation logic by successfully leveraging calculation groups
- When working with composite models using Power BI datasets or Azure Analysis Services and calculation groups, keep in mind wholesale vs non-wholesale and the interaction with calculation groups



Resources

Calculation groups

aka.ms/PowerBICalculationGroups

Different storage modes in Power BI

aka.ms/PowerBIStorageModes

Using Direct Query in Power BI

aka.ms/DirectQueryPBI

Guidance for Composite Models

aka.ms/CompositeModelsPBI

Relationship evaluation


aka.ms/DirectQueryRelationshipEvalPBI

DirectQuery for Power BI datasets and Azure Analysis Services (preview)

aka.ms/PBIDirectQueryDatasets

Model relationships in Power BI

aka.ms/ModelRelationshipsPBI



Want to review all this content?

Webinar Series: Mastering Data Modeling with Power BI

Episode 1 - Data Modeling 101: Increasing the Impact of Power BI

by Jeroen ter Heerdt and Marc Lelijveld

[Watch now](#)

Webinar Series: Mastering Data Modeling with Power BI

Episode 2 - Learn Advanced Data Modeling with Power BI

by Jeroen ter Heerdt and Marc Lelijveld

[Watch now](#)

Webinar Series: Mastering Data Modeling with Power BI

Episode 3 - Data Modeling for Experts with Power BI

by Jeroen ter Heerdt and Marc Lelijveld

[Watch now](#)

Webinar Series: Data Modeling with Power BI

Episode 4 - Calculation Groups and Composite Models

[Watch now](#)


Find them all here:

docs.microsoft.com/en-us/power-bi/fundamentals/webinars

LET'S
RECAP...



Thanks for attending!



Marc Lelijveld

Data & Analytics consultant
Macaw Netherlands

✉ Marc.Lelijveld@outlook.com

🐦 [@MarcLelijveld](https://twitter.com/MarcLelijveld)

in [linkedin.com/in/MarcLelijveld](https://www.linkedin.com/in/MarcLelijveld)

🌐 Data-Marc.com

